

Spark SQL motivation

192.168.0.38	WARNING	Something bad could happen
192.168.0.88	INFO	Just an info message passing by
192.168.0.5	WARNING	Something bad could happen
192.168.0.36	ERROR	When production fails in despair, whom are you gonna call?
192.168.0.27	INFO	Just an info message passing by

192.168.0.38	USA
192.168.0.88	RUSSIA
192.168.0.5	CHINA
192.168.0.36	USA
192.168.0.27	RUSSIA

Declarative

```
SELECT country, code FROM table1  
JOIN table2  
WHERE table1.ip = table2.ip  
AND table1.code != "INFO"
```

Imperative

```
rdd1 = sc.textFile("log.txt")
rdd2 = sc.textFile("ips.txt")
table1 = rdd1.map(lambda x: x.split("\t"))
table2 = rdd2.map(lambda x: x.split("\t"))
table1.cartesian(table2)
    .filter(lambda (x, y): x[0] == y[0])
    .filter(lambda (x, y): x[1] != "INFO")
    .map(lambda (x, y): (y[1], x[1]))
    .count()
```

Performance comparison

- 10,000,000 rows of logs (496 MB, 20 partitions)
- 100,000 rows of IPs (2 MB, 2 partitions)
- 10 executors, 2 cores & 4GB per executor

```
rdd1 = sc.textFile("log.txt")
rdd2 = sc.textFile("ips.txt")
table1 = rdd1.map(lambda x: x.split("\t"))
table2 = rdd2.map(lambda x: x.split("\t"))
table1.cartesian(table2)
    .filter(lambda (x, y): x[0] == y[0])
    .filter(lambda (x, y): x[1] != "INFO")
    .map(lambda (x, y): (y[1], x[1]))
    .count()
```


Details for Job 9

Status: RUNNING

Active Stages: 1

▶ Event Timeline

▼ DAG Visualization



Active Stages (1)

Stage Id ▾	Description		Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
13	count at <python-input-22-77368422c974>:2	+details (kill)	2017/07/14 19:34:05	2.0 h	<div><div></div>8/40</div>				

```
table1.createOrReplaceTempView("table1")  
table2.createOrReplaceTempView("table2")
```

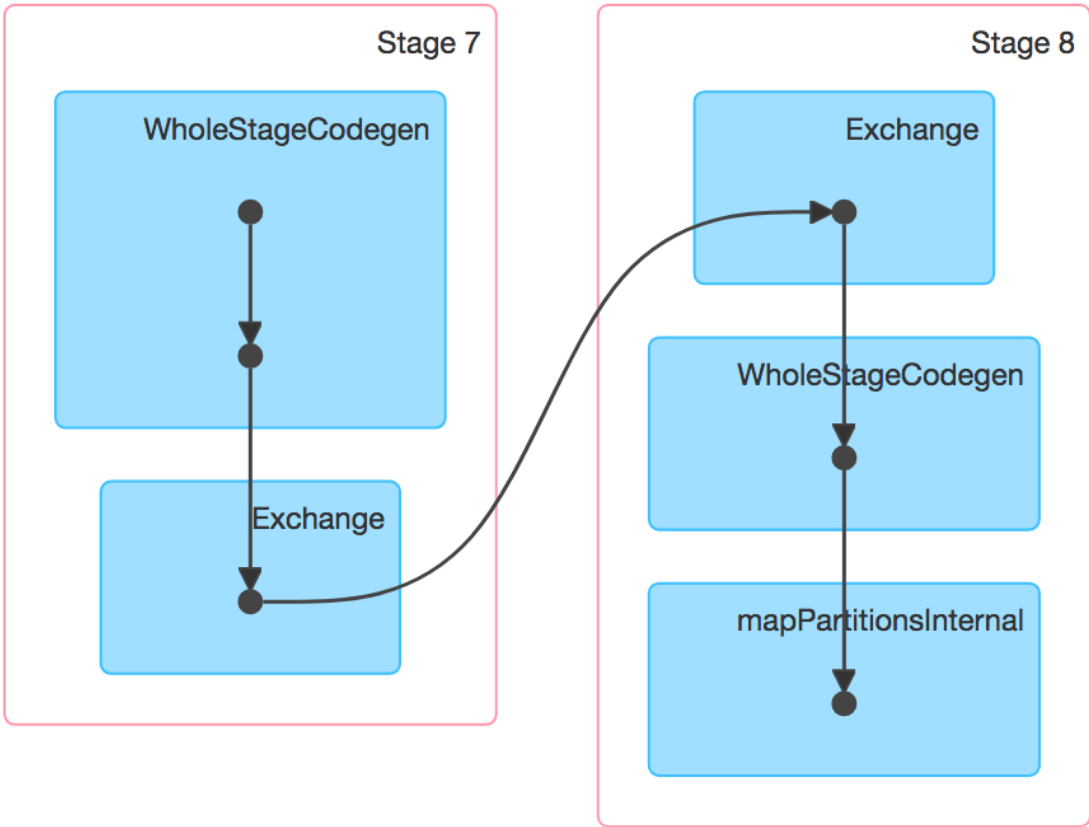
```
query = """  
    SELECT country, code FROM table1  
    JOIN table2  
    WHERE table1.ip = table2.ip  
    AND table1.code != INFO  
    """
```

```
spark.sql(query).count()
```

Details for Job 6

Status: SUCCEEDED
Completed Stages: 2

- ▶ Event Timeline
- ▼ DAG Visualization



6.6 seconds

Completed Stages (2)

Stage Id ▾	Description		Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
8	count at NativeMethodAccessorImpl.java:0	+details	2017/07/14 19:13:31	99 ms	1/1			1180.0 B	
7	count at NativeMethodAccessorImpl.java:0	+details	2017/07/14 19:13:27	4 s	20/20	496.0 MB			1180.0 B

RDD is a low level API

```
rdd1 = sc.textFile("log.txt")
rdd2 = sc.textFile("ips.txt")
table1 = rdd1.map(lambda x: x.split("\t"))
table2 = rdd2.map(lambda x: x.split("\t"))
table1.cartesian(table2)
    .filter(lambda (x, y): x[0] == y[0])
    .filter(lambda (x, y): x[1] != "INFO")
    .map(lambda (x, y): (y[1], x[1]))
    .count()
```


```
rdd1 = sc.textFile("log.txt")
rdd2 = sc.textFile("ips.txt")
table1 = rdd1.map(lambda x: x.split("\t"))
table2 = rdd2.map(lambda x: x.split("\t"))
table1.join(table2)
    .map(lambda (x, y): (y[1], y[0]))
    .filter(lambda (x, y): y != "INFO")
    .count()
```


Details for Job 0

Status: SUCCEEDED

Completed Stages: 2

- ▶ Event Timeline
- ▼ DAG Visualization



12.8 seconds

Completed Stages (2)

Stage Id ▾	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
1	count at <magic-timeit>:257 +details	2017/07/14 19:13:02	2 s	22/22			55.3 MB	
0	join at <magic-timeit>:257 +details	2017/07/14 19:12:52	11 s	22/22				55.3 MB

Spark knows nothing about
your data




```
SELECT country, code FROM table1  
JOIN table2  
WHERE table1.ip = table2.ip  
AND table1.code != "INFO"
```

```
table1 = sc.textFile("log.txt")
table2 = sc.textFile("ips.txt")
table1.join(table2)
    .filter(lambda (x, y): y[0] != "INFO")
    .map(lambda (x, y): (y[1], y[0]))
    .count()
```

some text file



code defines
data structure



Spark knows nothing about
your computations

```
SELECT country, code FROM table1  
JOIN table2
```

```
WHERE table1.ip = table2.ip  
AND table1.code != "INFO"
```

```
rdd1 = sc.textFile("log.txt")
rdd2 = sc.textFile("ips.txt")
table1.join(table2)
    .filter(lambda (x, y): y != "INFO")
    .map(lambda (x, y): (y[1], y[0]))
    .count()
```

Summary

- RDD is imperative
- RDD is low-level
- Using RDD computations are opaque
- Using RDD data is opaque